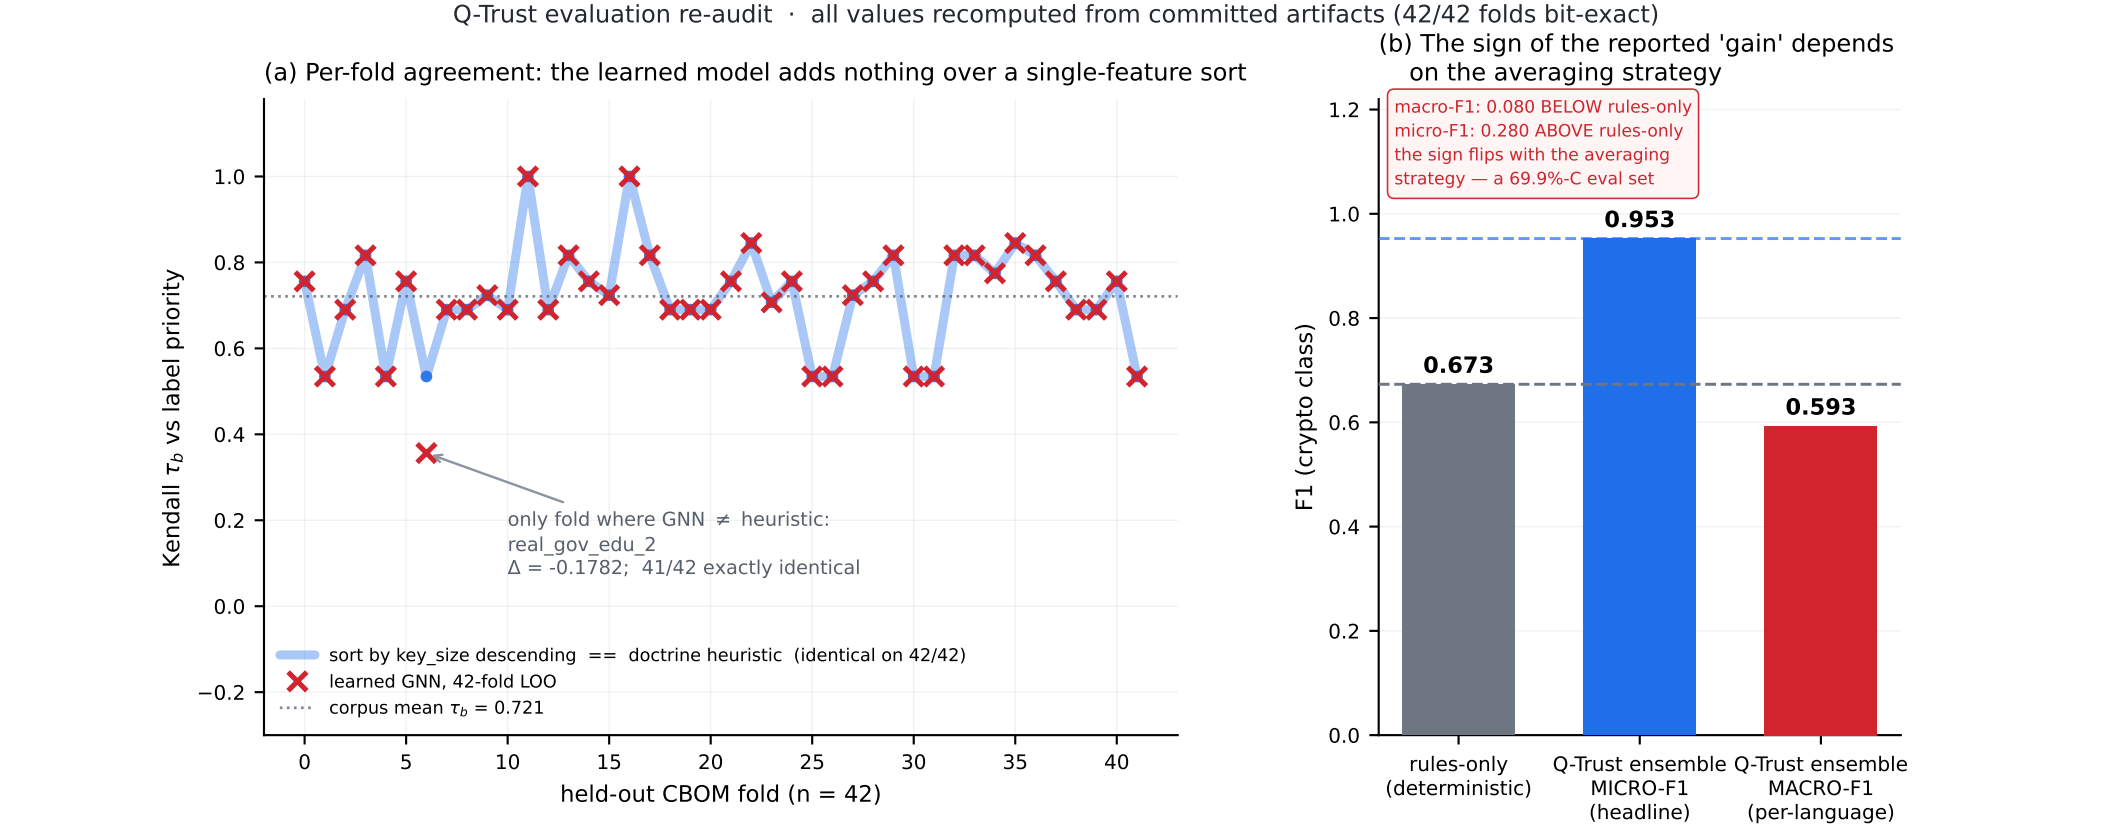
Q-Trust evaluation re-audit · all values recomputed from committed artifacts (42/42 folds bit-exact)
(a) Per-fold agreement: the learned model adds nothing over a single-feature sort
Kendall τ_b vs label priority
held-out CBOM fold (n = 42)
only fold where GNN ≠ heuristic:
real_gov_edu_2
Δ = -0.1782; 41/42 exactly identical
sort by key_size descending == doctrine heuristic (identical on 42/42)
learned GNN, 42-fold LOO
corpus mean τ_b = 0.721
(b) The sign of the reported 'gain' depends on the averaging strategy
F1 (crypto class)
macro-F1: 0.080 BELOW rules-only
micro-F1: 0.280 ABOVE rules-only
the sign flips with the averaging strategy — a 69.9%-C eval set
0.673
0.953
0.593
rules-only (deterministic)
Q-Trust ensemble MICRO-F1 (headline)
Q-Trust ensemble MACRO-F1 (per-language)